



# Project Documentation: Production-Ready BERT Sentiment Analysis API

This document provides in-depth technical details, architecture overview, setup guides, and troubleshooting information for the "**Production-Ready BERT Sentiment Analysis API**" project. It complements the README.md by offering a deeper dive into the project's internal workings.

## Contents

1. Project Overview .....	2
2. Project Architecture .....	2
3. Configuration Details .....	3
4. Local Setup & Installation .....	3
5. How to Run Components .....	5
6. API Reference.....	6
7. Codebase Details (Module by Module) .....	9
8. Common Issues & Troubleshooting .....	11
9. Future Enhancements .....	13
10. Contributing Guidelines .....	14
11. License & Acknowledgments .....	14

# 1. Project Overview

This project implements an end-to-end sentiment analysis solution for movie reviews using a fine-tuned **BERT** model. It demonstrates a complete machine learning pipeline from data acquisition and preprocessing to model training, evaluation, and deployment as a real-time **RESTful API** using **FastAPI**.

For a concise overview, key features, and initial setup instructions, please refer to the main README.md file.

## 2. Project Architecture

### 2.1. High-Level Flow

The project follows a standard machine learning pipeline:

- **Data Processing:** Raw IMDB movie review data is loaded, cleaned, and split into training, validation, and test sets.
- **Model Building:** A **BERT**-based classifier (using Hugging Face's TFBertModel) is constructed.
- **Training & Evaluation:** The model is fine-tuned on the training data, monitored using callbacks, and comprehensively evaluated on the test set.
- **API Deployment:** The trained model is served via a **FastAPI** application, allowing real-time sentiment predictions through HTTP requests.

### 2.2. Module Breakdown

The project is structured into modular Python files, each with a specific responsibility:

- `config.py`: Centralizes all configuration settings.
- `data_loader.py`: Handles IMDB dataset operations (loading, cleaning, splitting).
- `model.py`: Defines the **BERT** model's architecture.
- `trainer.py`: Manages the model training pipeline.
- `evaluator.py`: Conducts model evaluation and visualization.
- `app.py`: Implements the **FastAPI** inference **API**.
- `main.py`: Orchestrates the entire training and evaluation workflow.

- `test_api.py`: A separate script for programmatically testing the local **API**.
- `requirements.txt`: Lists all project dependencies.
- `.gitignore`: Specifies files/folders to be ignored by Git.

### 3. Configuration Details

All key parameters and paths are managed centrally in `config.py` using dataclasses for clarity and easy modification.

- **ModelConfig**: Defines model-specific hyperparameters (e.g., `model_name`, `max_length`, `learning_rate`).
- **TrainingConfig**: Controls training strategies and callbacks (e.g., `epochs`, `save_strategy`, `metric_for_best_model`).
- **ProjectConfig**: Manages project directory structure and file paths (e.g., `models_dir`, `outputs_dir`).

### 4. Local Setup & Installation

Follow these steps to set up and run the project on your local machine.

#### Prerequisites:

- **Python 3.11**: Recommended for stable **TensorFlow/Text** compatibility.
- **Git**: For cloning the repository.
- **Visual Studio Code** (or any IDE): For code editing and terminal access.
- **Microsoft C++ Build Tools** (for Windows users): Essential for compiling certain Python packages that contain C/C++ extensions. Download from [Microsoft Visual C++ Build Tools](#) (select "Desktop development with C++" workload during installation).

#### Installation Steps:

1. **Clone the repository:**
2. `git clone https://github.com/EmadAliEmad/BERT-Fine_tuning-for-Movie-Sentiment-Analysis.git`
3. `cd BERT-Fine_tuning-for-Movie-Sentiment-Analysis`

#### 4. Download the Trained Model (best\_model.h5):

The trained model is not included in this repository due to its large size. It must be downloaded separately.

- Go to the Kaggle Notebook output associated with this project: [Kaggle Notebook Output \(Version with Trained Model\)](#)
- Download best\_model.h5 from the models/ folder within that specific version's output.
- Create a models/ directory in your local project root: `mkdir models`
- Place the downloaded best\_model.h5 file inside the models/ directory.

#### 5. Create and activate a Python virtual environment:

It's highly recommended to use a virtual environment to manage project dependencies isolation.

6. `# Using Python 3.11 directly (if installed as 'python3.11' or 'py -3.11')`
7. `py -3.11 -m venv venv_api`
8. `# Or simply 'python -m venv venv_api' if 3.11 is your default Python.`
- 9.
10. `# Activate the virtual environment:`
11. `.\venv_api\Scripts\activate # On Windows (Command Prompt or PowerShell)`
12. `source venv_api/bin/activate # On macOS/Linux/Git Bash`

IGNORE\_WHEN\_COPYING\_START

content\_copy download

Use code [with caution](#). Bash

IGNORE\_WHEN\_COPYING\_END

#### 13. Install the required dependencies:

All project dependencies are listed in requirements.txt.

14. `pip install -r requirements.txt`

IGNORE\_WHEN\_COPYING\_START

content\_copy download

Use code [with caution](#). Bash

IGNORE\_WHEN\_COPYING\_END

(You may see WARNING messages about dependency conflicts; these can usually be ignored as long as no ERROR prevents installation. For local Windows setup, ensure C++ Build Tools are installed if compilation errors occur.)

## 5. How to Run Components

### 5.1. Training & Evaluation Pipeline (main.py)

The main.py script orchestrates the entire training and evaluation process. It's designed to either load an existing model or train a new one if not found.

To run the full pipeline (training if model not found, then evaluation):

```
python main.py
```

IGNORE\_WHEN\_COPYING\_START

content\_copy download

Use code [with caution](#). Bash

IGNORE\_WHEN\_COPYING\_END

Output: This will log detailed progress to the console (using **Rich**) and save logs to logs/app.log, training history to outputs/training\_history.json, and generate interactive HTML plots in outputs/.

If best\_model.h5 is not in models/, this will trigger a full training run (which is resource-intensive and requires a **GPU** or significant **CPU** time).

### 5.2. FastAPI Inference API (app.py)

The app.py script runs the **FastAPI** server for real-time sentiment predictions.

To run the **API** locally:

```
uvicorn app:app --reload --host 0.0.0.0 --port 8000
```

IGNORE\_WHEN\_COPYING\_START

content\_copy download

Use code [with caution](#). Bash

IGNORE\_WHEN\_COPYING\_END

- `uvicorn app:app`: Tells **Uvicorn** to run the app instance (our **FastAPI** application) from the `app.py` file.
- `--reload`: Restarts the server automatically when code changes are detected (useful for development).
- `--host 0.0.0.0`: Makes the server accessible from all network interfaces (allows external connections).
- `--port 8000`: Specifies the port the **API** will listen on.

Output: The **API** will be running on `http://127.0.0.1:8000`. Keep this terminal open while using the **API**.

**Note:** If running on a remote server/notebook (like Kaggle), you would typically use **ngrok** (as demonstrated in the Kaggle Notebook) to expose this local port to a public URL.

## 6. API Reference

The **FastAPI** application automatically generates interactive **API** documentation.

### 6.1. Base URL

- Local: `http://127.0.0.1:8000`
- Documentation (Swagger UI): `http://127.0.0.1:8000/docs`

### 6.2. Endpoint: GET /health

- **Purpose:** Provides a simple health check for the **API**. It verifies if the server is running and if the **BERT** model and tokenizer have been successfully loaded into memory.
- **Request:**

- **Method:** GET
- **URL:** /health
- **Parameters:** None
- **Response:**
  - **Status Code:** 200 OK
  - **Body (JSON):**
    - {
    - "status": "ok",
    - "model\_loaded": true,
    - "tokenizer\_loaded": true
    - }

IGNORE\_WHEN\_COPYING\_START

content\_copy download

Use code [with caution](#). Json

IGNORE\_WHEN\_COPYING\_END

### 6.3. Endpoint: POST /predict

- **Purpose:** Accepts text input(s) and returns sentiment prediction(s) (Positive/Negative) along with a confidence score.
- **Request:**
  - **Method:** POST
  - **URL:** /predict
  - **Content-Type:** application/json
  - **Body (JSON - adheres to PredictionRequest schema):**
    - {
    - "texts": [

- "This movie was an absolute masterpiece! I loved every moment of it.",
- "The plot was confusing and the acting was terrible. A complete waste of my time and money."
- ]
- }

IGNORE\_WHEN\_COPYING\_START

content\_copy download

Use code [with caution](#). Json

IGNORE\_WHEN\_COPYING\_END

- **Response:**

- **Status Code:** 200 OK
- **Body (JSON - adheres to PredictionResponse schema):**
- {
- "predictions": [
- {
- "text": "This movie was an absolute masterpiece! I loved every moment of it.",
- "sentiment": "Positive",
- "confidence": 0.998
- },
- {
- "text": "The plot was confusing and the acting was terrible. A complete waste of my time and money.",
- "sentiment": "Negative",
- "confidence": 0.995
- }



- ]
- }

IGNORE\_WHEN\_COPYING\_START

content\_copy download

Use code [with caution](#). Json

IGNORE\_WHEN\_COPYING\_END

- **Error Status Code:** 422 Unprocessable Entity (if request body format is incorrect).

## 7. Codebase Details (Module by Module)

Here's a detailed description of each Python module in the project:

- `app.py`:
  - **Role:** The core **FastAPI** application script.
  - **Key components:** Defines **API** endpoints (`/health`, `/predict`), handles model/tokenizer loading on startup (`@app.on_event("startup")`), and contains the inference logic for sentiment prediction. It includes `PredictionRequest` and `PredictionResponse` **Pydantic** models.
- `config.py`:
  - **Role:** Centralized configuration management.
  - **Key components:** Contains `ModelConfig` (model hyperparameters), `TrainingConfig` (training strategies), and `ProjectConfig` (file paths and directory management). Uses dataclasses for clean structure.
- `data_loader.py`:
  - **Role:** Data processing pipeline.
  - **Key components:** `DataProcessor` class with methods for `load_imdb_dataset` (decoding IMDB integer sequences to text), `clean_text` (standardizing text preprocessing), and `create_data_splits` (generating train/val/test sets).
- `evaluator.py`:

- **Role:** Model evaluation and visualization.
  - **Key components:** ModelEvaluator class provides predict method (for test data), evaluate\_model (calculates metrics like accuracy, F1-score, ROC AUC, and generates classification report/confusion matrix), and plotting functions (plot\_training\_history, plot\_confusion\_matrix, plot\_roc\_curve) using **Plotly** for interactive HTML outputs.
- logger.py:
  - **Role:** Professional logging setup.
  - **Key components:** setup\_logging function configures **Loguru** with **RichHandler** for visually appealing console output and file logging (e.g., logs/app.log), ensuring all project activities are monitored.
- main.py:
  - **Role:** The main script orchestrating the full project pipeline.
  - **Key components:** Calls setup\_logging, print\_project\_info, and controls the sequence of data loading, model building/loading (checking for existing model on disk), training, and evaluation. It links all other modules to execute the end-to-end workflow.
- model.py:
  - **Role:** Defines the **BERT** model architecture.
  - **Key components:** BERTSentimentClassifier (a custom tf.keras.Model subclass) which loads TFBertModel from **transformers** and adds a classification head. BERTModelBuilder provides a static method build\_functional\_model for constructing the model using **Keras Functional API**.
- trainer.py:
  - **Role:** Encapsulates the model training pipeline.
  - **Key components:** BERTTrainer class handles compile\_model (optimizer, loss, metrics setup), setup\_callbacks (e.g., ModelCheckpoint, EarlyStopping, ReduceLROnPlateau, TensorBoard, CSVLogger), and the train method (fitting the model on data). It also manages save\_training\_artifacts for history and configurations.

- requirements.txt:
  - **Role:** Lists all Python library dependencies with specific versioning for exact environment replication.
- test\_api.py:
  - **Role:** A standalone Python script for programmatically testing the locally running **FastAPI API** endpoints (health check and prediction).
- .gitignore:
  - **Role:** Specifies files and directories that Git should ignore and not include in the version control history (e.g., large model files, virtual environments, log files).

## 8. Common Issues & Troubleshooting

This section details common problems encountered during this project's development and their solutions.

- **ModuleNotFoundError: No module named 'xyz':**
  - **Reason:** The required Python library (xyz) is not installed in the active virtual environment.
  - **Solution:** Run `pip install -r requirements.txt` after activating your virtual environment. Ensure `!pip install pyngrok -q` is run directly in the Kaggle cell that imports it.
- **TypeError: Object of type float32 is not JSON serializable:**
  - **Reason:** TensorFlow/NumPy float32 data types are not natively supported by the standard JSON format.
  - **Solution:** Convert float32 values to standard Python float before JSON serialization (implemented in `trainer.py` in `save_training_artifacts`).
- **huggingface\_hub.utils.\_validators.HFValidationError: Repo id must use alphanumeric chars...:**
  - **Reason:** Attempting to re-initialize a tokenizer with an already-loaded tokenizer object itself, instead of its string name.

- **Solution:** Ensure the tokenizer is passed directly as an initialized object where needed, or with its string name when initializing from scratch (fixed in evaluator.py).
- **NameError: name 'tf' is not defined:**
  - **Reason:** The **tensorflow** library was not imported (e.g., import tensorflow as tf) in the specific .py file where tf was being used.
  - **Solution:** Add import tensorflow as tf at the beginning of the relevant module (fixed in evaluator.py).
- **ValueError: Unknown layer: 'TFBertModel'. Please ensure you are using a keras.utils.custom\_object\_scope...:**
  - **Reason:** **Keras** does not natively recognize TFBertModel (which comes from Hugging Face **transformers**) when loading a saved model.
  - **Solution:** Pass custom\_objects={'TFBertModel': TFBertModel} as an argument to tf.keras.models.load\_model() when loading the best\_model.h5 file (fixed in main.py and app.py).
- **ValueError: Input 0 of layer "bert\_sentiment\_classifier" is incompatible with the layer: expected shape=(None, 128), found shape=(None, X):**
  - **Reason:** The tokenizer was not padding input sequences to the exact max\_length (128) expected by the **BERT** model.
  - **Solution:** Explicitly set padding='max\_length' in the tokenizer's call (fixed in app.py).
- **Health Check Status: 404 or JSONDecodeError during API testing:**
  - **Reason:** This often indicates the **API** is not reachable at the specified **URL**. Common causes include:
    - **Ngrok URL Mismatch:** The ngrok\_url in the testing script/browser is outdated.
    - **Ngrok Service Down:** The **ngrok** tunnel itself is not active or crashed.
    - **Uvicorn Server Down:** The **FastAPI** application (**Uvicorn**) crashed or is not running.
  - **Solution:**

- Always copy the NEW **Ngrok Tunnel URL** from the output of the **API** startup cell (Cell 14) and paste it into the testing cell (Cell 15).
  - Ensure no other **ngrok** sessions are running (check **ngrok** dashboard and kill any active sessions).
  - Ensure the **API** startup cell (Cell 14) executes successfully and continuously.
- **The ngrok process errored on start: authentication failed: Your account is limited to 1 simultaneous ngrok agent sessions. (ERR\_NGROK\_108):**
    - **Reason:** Your **ngrok** free account only allows one active tunnel at a time. A previous session (from another notebook, local machine, or a crashed Kaggle session) might still be active.
    - **Solution:** Manually stop all active **ngrok** sessions from your **ngrok** dashboard. The `ngrok.kill()` command in our setup cells also helps in cleaning up residual processes within the Kaggle environment.

## 9. Future Enhancements

The project provides a solid foundation. Here are potential areas for further development:

### Web User Interface (Frontend):

- Integrate a user-friendly web interface using **Streamlit** (already in requirements.txt), **Flask**, or **Dash** to allow users to input text and see real-time sentiment predictions.

### Model Optimization & Fine-tuning:

- Experiment with different **BERT** variants (e.g., bert-large-uncased for potentially higher accuracy, but more resources) or other **Transformer** models.
- Explore advanced fine-tuning techniques (e.g., learning rate schedulers, weight decay).
- Fine-tune on a domain-specific dataset if targeting a particular industry (e.g., financial sentiment).

### Advanced NLP Features:

- Implement more sophisticated text preprocessing (e.g., handling slang, emojis).

- Expand to emotion detection (e.g., joy, sadness, anger) or aspect-based sentiment analysis.
- Support multi-lingual sentiment analysis by using multi-lingual **BERT** models.

#### Deployment & MLOps:

- Containerize the application using **Docker** for easier packaging and deployment across different environments.
- Deploy the **API** to a permanent cloud platform (e.g., Hugging Face Spaces, **AWS Lambda**, **Google Cloud Run**, **Azure App Service**) for continuous availability.
- Add **API** authentication/authorization, rate limiting, and more detailed **API** logging for production readiness.

#### Monitoring & Alerting:

- Implement model performance monitoring and data drift detection.

## 10. Contributing Guidelines

We welcome contributions to this project!

1. Fork the repository.
2. Clone your forked repository: `git clone https://github.com/YourUsername/BERT-Fine_tuning-for-Movie-Sentiment-Analysis.git`
3. Create a new branch for your feature or bug fix: `git checkout -b feature/your-feature-name`
4. Make your changes, add comments, and ensure code quality.
5. Test your changes thoroughly.
6. Commit your changes: `git commit -m "feat: Add new feature X"` (use conventional commits).
7. Push to your branch: `git push origin feature/your-feature-name`
8. Open a Pull Request to the main branch of the original repository.

## 11. License & Acknowledgments

This project is licensed under the **MIT License**.

**Acknowledgments:**

- **Kaggle:** For providing the **GPU**-enabled notebook environment.
- **Hugging Face:** For the excellent **Transformers** library and pre-trained **BERT** models.
- **TensorFlow:** For the deep learning framework.
- **FastAPI:** For the fast and modern web framework.
- **Rich & Loguru:** For enhancing logging and console output.